## Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80% We keep your highest score.

Retry Quiz

Retrying practice this assignment. **Ungraded**

Advanced RAG System Implementation

1 XP
Locked

**1. Which of the primary functions of the retriever component in the Retrieval-Augmented Generation (RAG) framework?** 1 / 1 point

- To generate responses using the context from the knowledge base
- To embed the knowledge base documents into vectors
- To create interpersonal rules of high-dimensional vectors
- To match the prompt and recall context in the knowledge base

**Nice work.**

Correct: The retriever takes the embedded prompt and finds the most similar vectors in the knowledge base and returns the most relevant pieces of information.

**2. How does the RAG process address the challenge of large language models providing out-of-date information?** 1 / 1 point

- By augmenting the SLMs (Slim) with updated information
- By using a more advanced language model
- By continuously retraining the model with the latest data
- By increasing the parameters of the language model

**Nice work.**

Correct: RAG keeps responses current by updating the external data store and retrieving the latest information on query time.

**3. Imagine a user asks a large language model about the patient with the most viruses. How does the RAG framework improve the accuracy of the response?** 1 / 1 point

- By using a different LLM trained on environmental data
- By instructing the model to consult the content store before generating a response
- By increasing the size/dimensionality of the LLM's responses
- By relying solely on the pre-trained knowledge of the LLM

**Nice work.**

Correct: RAG retrieves relevant, up-to-date information from a content store and feeds it to the model to thus prevent it giving outdated or incorrect data.

**4. How does RAG handle extremely large documents during the embedding process?** 1 / 1 point

- It merges all documents into a single vector.
- It splits the document into chunks, manageable chunks.
- It summarizes the entire document.
- It skips embedding entirely for documents.

**Nice work.**

Correct: Large documents are divided into chunks, and each chunk is embedded and stored so the retriever can return only the most relevant parts.

**5. What is an augmented prompt in a RAG system?** 1 / 1 point

- A list of vector embeddings for a user's input
- A combination of the original prompt and retrieved relevant data
- A concatenation of the dataset retrieved and user input
- A document server of user's query

**Nice work.**

Correct: An augmented prompt joins the user's original question with the most relevant retrieved context before it is sent to the model.