coursera | IBM

Vector Databases for RAG: An Introduction

Today's goals

- Complete any learning items - 0/4
- Complete a module
- Complete a practice test

Today's Skill Points 11 SP

Introduction

Introduction to Vector Databases and Chroma DB

Welcome to the Course

- IBM Full-Stack Capstone Intro
 Reading • 5 min
- Video Course Introduction
 Video • 3 min
- Course Overview
 Reading • 5 min
- Hot and Agile: AI-Infused Local Food Hub
 Video • 3 min
- Reading: Helpful Tips for Course Completion
 Graded Peer Task • 5 min

Introduction to Vector Databases and Similarity Search

- Vector Database Concepts
 Video • 10 min
- Reading: From Database Rows: Traditional Evolution
 Graded Peer Task • 15 min
- Vector Database Types
 Video • 10 min
- Applications of Vector Databases
 Video • 10 min
- Reading: Similarity Search
 Graded Peer Task • 8 min
- Lab: Similarity Search by Hand
 Graded Quiz Task • 30 min
- (Optional) Interactive Lesson Recap: Podcast ON Demand
 Interactive Item Task • 15 min
- Practice Quiz: Introduction to Vector Databases
 Practice Assignment • 12 min

Chatting! Chroma DB

- Chroma DB Key Concepts and Architecture
 Video • 10 min
- Reading: Chroma and Metadata
 Graded Peer Task • 8 min
- Reading: Similarity Search and Metadata Filtering
 Graded Peer Task • 8 min
- Lab: Similarity Search in Action Using Chroma DB and Python
 Graded Quiz Task • 30 min
- (Optional) Interactive Lesson Recap: Podcast ON Demand
 Interactive Item Task • 15 min
- Practice Quiz: Exploring Chroma DB
 Practice Assignment • 12 min

Module Summary and Evaluation

- Reading: Summary and Highlights: Introduction to Vector Databases and Chroma DB
 Reading • 5 min
- Reading: Next Steps: Introduction to Vector Databases and Chroma DB
 Graded Peer Task • 8 min
- Graded Quiz: Introduction to Vector Databases and Chroma DB
 Graded Quiz • 30 min • 4 min
- (Optional) Discuss in Group: Meet and Greet
 Discussion Prompt • 10 min

Module 2

Vector Databases for Recommendation Systems and RAG

Course Overview

Shift Left and Peer Assist Opportunities ... Shift Left and Innovation Development and Strategy

0:00 / 0:00

Hello, and welcome to the **Vector Databases for Web AI Introduction** course.

Gain expertise in using vector databases and improve your data retrieval skills with hands-on practice! Start by exploring the fundamentals of similarity search and cluster vector databases, understanding how they differ from traditional databases, and recognizing their importance in recommendation systems and Retrieval Augmented Generation (RAG) applications.

Build on this foundation with practical experience using Chroma DB—one of the leading vector database solutions. Learn how to create, update, and delete collections, and perform efficient similarity searches. You'll also work with datasets, understand key concepts such as vector operations and relations in Chroma to develop a strong grasp of Chroma DB's functionality.

Apply what you've learned by creating a real-world recommendation system powered by Chroma DB and an interesting model from Hugging Face. This project will enhance your understanding of how vector databases improve data search and result in AI-driven applications.

Throughout the course, engage in practical exercises to refine your skills in database management, similarity search techniques, and recommendation system design. You'll also explore the challenges and limitations of vector databases and similarity searches, which are essential components of this field of study.

This course is part of the [IBM Full-Stack AI Engineering Specialization](). By completing the program, you'll understand how to effectively leverage vector databases and traditional databases based on their functionality and use cases. You will be able to perform fundamental database operations in Chroma DB, including updating, deleting, and managing collections. Additionally, you'll gain hands-on experience with similarity search, including the nuances between white-box, round-robin, brute-force, and similarity-level thresholds and apply them to classify and sort documents, and you will be able to develop recommendation systems using these techniques.

Prerequisites

This course is accessible to a wide range of professionals, but having some knowledge and experience of the Python programming language is helpful. Familiarity with software applications and machine learning models and some familiarity with databases is advantageous but not required.

Objectives

After completing this course, you will be able to:

- Differentiate between vector databases and traditional databases based on their functionality and use cases.
- Execute fundamental database operations in Chroma DB, including updating, deleting, and managing collections.
- Understand and apply similarity search techniques, both manually and with Chroma DB, to retrieve relevant information from text and multimodal, and develop recommendation systems using Chroma DB.
- Develop a thorough, hand-on grasp on understanding of key internal mechanisms within RAG.

Course Outline

This course consists of 2 learn modules.

Module 1: Introduction to Vector Databases and Chroma DB

- Lesson 1: Welcome
- Lesson 2: Introduction to Vector Databases and Similarity Search
- Lesson 3: Exploring Chroma DB

Key Topics:

- Concept of Vector Databases
- Comparing vector databases with traditional databases
- Types of vector databases
- Use cases and applications for vector databases
- Similarity search
- Concepts and intricacies of Chroma DB
- Chroma DB filtering
- Similarity search with Chroma DB

Module 2: Vector Databases for Recommendation Systems and RAG

- Lesson 1: Chroma DB: Database Operations
- Lesson 2: Develop a Recommendation System and Connect (Layered Concepts to RAG)

Key Topics:

- Essential Chroma DB database operations
- Reducing data with similarity search
- How vector databases power RAG-applications
- Building a recommendation system with Chroma DB

Congratulations on completing this course. Boost your career with systems analysis skills. We hope you enjoy your learning journey.

Mark as completed

Like Dislike Report an issue

Go to next item