Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 1 to remove your highest score.

Retry Graded Assignment

Applied Generative AI Development and Strategy

1. **How does a vector database facilitate the implementation of a recommendation system?** 1 / 1 point

   - By optimizing SQL queries for faster data retrieval.
   - By storing and retrieving embeddings that represent item features.
   - By generating new data points based on existing user data.
   - By directly storing user preferences as textdata.

   Nice work

   A vector database stores embeddings which are numeric representations features, enabling similarity searches essential for recommendation systems.

2. **Which is a key advantage of using a vector database for Retrieval Augmented Generation (RAG)?** 1 / 1 point

   - Frequent real-time data inserting for large datasets.
   - Hardware for efficient similarity searches using vector embeddings.
   - Guaranteed data freshness in augmented prompts from the original user prompt and the retrieved text.
   - Specialized advanced visualization tools for data analysis.

   Nice work

   Vector databases enhance RAG by enabling efficient similarity searches, which are crucial for retrieving relevant information using vector embeddings.

3. **In the context of vector databases, how can similarity search be applied to improve user experience in a recommendation system?** 1 / 1 point

   - By identifying items that are closely related in feature space to the user's preferences.
   - By using SQL queries to filter out irrelevant data.
   - By generating random recommendations based on user history.
   - By storing user data in a structured tabular format.

   Nice work

   Similarity searches in vector databases identify items that are closely related in feature space, which helps in providing personalized recommendations.

4. **How do embedding models and vector databases work together in a recommendation system?** 1 / 1 point

   - Embedding models replace vector databases to simplify data storage.
   - Embedding models visualize data stored in vector databases.
   - Vector databases always store iconic embeddings for recommendations.
   - Vector databases store embeddings that are used for similarity searches.

   Nice work

   Embedding model generates vectors that are stored in vector databases, facilitating similarity searches for recommendations.

5. **Which of the following can you not do with the 'modify' method in Chroma DB?** 1 / 1 point

   - Change the distance metric.
   - Change the collection metadata.
   - Change the collection name.
   - Add a key-value pair to the collection metadata.

   Nice work

   In Chroma DB, the distance metric must be specified when the collection is created. To change it later, you need to store the entire data using the 'modify' method will not work.

6. **What is a commonly effective chunking strategy?** 1 / 1 point

   - Choosing a chunk size that is large enough to keep the meaning clear.
   - Re-embedding data after changes.
   - Using a different embedding model for query and document embedding.
   - Using vector database for storage.

   Nice work

   Using a different embedding model for query and document embedding will typically lead to poor retrieval performance and is a critical ‘do not do’.

7. **Which of the following RAG tasks is typically handled outside the vector database?** 1 / 1 point

   - Chunking
   - Embedding prompts
   - Storing embeddings
   - Retrieving relevant chunks

   Nice work

   Chunking is typically performed outside the vector database, with the resulting documents/chunks then passed to the vector database for embedding and storage.