

AI Fundamentals

Module 1

Collaborate with AI

Module 2

Practice using AI

Module 3

Learn how AI works

Module 4

Design great prompts

Module 5

Level up your prompts

● Prompt chaining strategies

Reading = 5 min

● Practice prompt chaining in Gemini

Lab = 10 min

● Make what's ahead for giving Gemini context in a new chat

Video = 2 min

● Generate clouds with AI

Reading = 4 min

● Create custom images with Gemini

Lab = 10 min

Module 6

Use AI responsibly

● Responsible AI: practice

Video = 4 min

● What, what should you do if you get a bad response from Gemini?

Video = 4 min

● Responsible AI tool practices

Reading = 5 min

● Learn about AI agents

Reading = 8 min

● Course challenge

Graded assignment = 10 min

● Claim your course completion badge

Reading = 2 min

● Get started on the next course

Reading = 2 min

Learn about AI agents

Some types of AI can do more than execute a single task. They can autonomously handle complex, multistep workflows with your guidance and oversight, helping you achieve broad goals. Unlike standard AI chatbots, you can choose to let these AI agents access the tools you use every day, like email, calendars, or documents, and take action on your behalf and under your control.

What can AI agents do for you?

Think of an agent as an AI system that you can delegate complex, multistep tasks to. This skill helps free you so you can focus on activities that need your expertise, like strategic thinking.

Here are a few examples of how you might use an AI Agent:

- **Conduct market research:** You can use an agent to analyze new competitors, find their key features, and synthesize recent customer reviews into a weekly team email.
- **Streamline sales and CRM:** You can set up an agent to go through your inbox, identify new sales leads, extract their contact information, and automatically create new entries for them in your CRM.
- **Assist with forecasting demand:** You can have an agent analyze sales and inventory data, predict future demand for your products, and order the necessary supplies to make sure you're prepared.

Understand how AI agents work

An AI agent operates using three core components:

1. **Models:** Models are computer programs trained on vast amounts of data, which allows them to recognize patterns. When you give an AI agent a goal, the model uses its training to predict which tools are necessary to achieve that goal and how to use the tools available to complete it.
2. **Permissions:** AI agents connect to other software and tools through specific permissions you grant. These settings set guardrails, controlling how the agent executes tasks within the parameters you set. For example, if you assign an agent to manage your inbox, you can restrict its access to only read email messages.
3. **Goals:** A goal is the broad objective you provide an AI agent, such as "Organize the Q3 kickoff meeting." Once a goal is set, the agent operates in a continuous loop until the goal is complete.
 - **Analyze:** The agent uses its model to determine the necessary first task, like connecting with your calendar to find free time.
 - **Act:** Once permissions are granted, the agent executes that task, like finding a time for the team to connect.
 - **Observe:** The agent reviews the result of its action and ensures it's aligned with the goal, like checking the selected time to make sure everyone is available.

After completing the Observe step, the agent begins the loop again by moving to the next task, like drafting and sending the calendar invitations. The agent continues this loop until the meeting is successfully scheduled.

Your guide to working with agents responsibly

As AI agents become more capable, your guidance and judgment become even more important. When deciding whether to delegate a goal to an agent, keep these two principles in mind:

1. **Risk:** What's the impact of a mistake? For example, drafting an internal memo may be low risk. But using an agent to make purchases on your behalf might be high risk depending on the size of the transaction.
2. **Reversibility:** How easily can an action be undone? Drafting an email and saving it to your drafts folder is highly reversible. Posting directly to a public social media account is not.

For high-risk or irreversible goals, you should maintain tighter control through a human-in-the-loop approach. Set clear boundaries, define acceptable parameters, and maintain oversight at each critical step.

The best way to use AI agents responsibly is to start small. Look for agentic features in the software and tools you're already using, and try giving an agent a simple, low-risk goal.

Best practices for directing AI agents

As you start to use AI agents throughout your work, consider the following best practices:

- **Start with a clear, specific goal:** Provide a precise description of what you're trying to accomplish. "Plan a marketing campaign" is vague. "Draft a 3-email sequence for our new product launch" is clear and actionable.
- **Manage permissions wisely:** Be mindful of what you approve. If an agent needs to read emails to find contacts, grant only that specific permission. Avoid giving master permissions that grant blanket access to your entire app.
- **Always evaluate:** You're responsible for any work the agent produces. Always check the work, especially on high-stakes tasks.
- **Iterate and refine:** Treat the agent as a collaborator. If the first result isn't right, refine your instructions or provide more context.

AI agents can be powerful partners, but your success depends on your ability to supervise them with clarity, judgment, and a strong vision.

Mark as completed